


Interview synthesis and your **no-slop** setup

From what users said to a setup that is ready to build

Class 4 · Friday, October 2, 2026 · 1:10 to 3:40 PM · Pupin 420

What we're doing today

No break today.



The rest of the semester, 1 of 3

Semester arc

Date	Class	Topic	Homework due
Today	Class 4	Synthesis and no-slop setup	HW 3: design work, interviews
Oct 9	Class 5	Build kickoff and landing pages	HW 4: synthesis, personas, setup
Oct 16	Class 6	Driving traffic	HW 5: landing page, demand test
Oct 23	Class 7	MVP demos	HW 6: working MVP

The rest of the semester, 2 of 3

Semester arc

Date	Class	Topic	Homework due
Oct 30	Class 8	What to build next	HW 7: MVP learnings
Nov 6	Class 9	Get users and analytics	HW 8: roadmap and CLAUDE.md
Nov 13	Class 10	Cold start	HW 9: growth plan, analytics
Nov 20	Class 11	Showcase and fundraising	HW 10: cold start strategy

The rest of the semester, 3 of 3

Semester arc

Date	Class	Topic	Homework due
Nov 27	No class	Thanksgiving	
Dec 4	Class 12	Feature freeze and pitching	HW 11: iterate, then freeze
Dec 11	Class 13	Unit economics, deck workshop	HW 12: pitch deck draft 1
Dec 18	Demo Day	Final pitches and live demos	HW 13: final deck

What you work on next, and why

Semester arc

Next class, Oct 9

- Class 5: Build kickoff and landing pages
- Homework 4: synthesis, personas, setup
- Why it matters: Your setup is done, so now you build. A landing page tests demand while the app takes shape.

After that, Oct 16

- Class 6: Driving traffic
- Homework 5: landing page, demand test
- Why it matters: Your copy is a guess until strangers see it. Traffic to a live page shows real demand.

AI news this week

News

This week in AI

AI news

Which one do you want to talk about?

Jev

A small model that returns a decision with a probability, not text. Early access, \$0.042 per million input tokens.

[TypeSafe docs](#)

Decision models

Seven more have followed since Jev launched, from OpenAI, Cloudflare, Perplexity, and others. Some are open weights.

[Cloudflare, Oct 1](#)

Sonnet 5.5

Out Sep 28, \$2 and \$10 per million tokens, same as Sonnet 5. Anthropic claims faster and cheaper per task.

[Anthropic, Sep 28](#)

A model that returns decisions

TypeSafe calls it a System One model: fast and narrow, not chatty.

What it is

- Send text and a question. Get a typed answer and a probability.
- **\$0.042** per million input tokens. Output is free.
- 70 to 500 milliseconds per call. Text only.

What you could build

- **Route** tickets to the right team.
- **Rank** leads against your own rubric.
- **Verify** a claim against its source.
- **Guard** against unsafe or off-topic messages.

Decision models

Big names followed Jev

Same idea: a typed answer and a probability, no prose.

Model	Maker	Date	Price or access
Decisions API	OpenAI	Sep 29	Limited preview, price not stated
Clef, Clef-flash	Cloudflare	Oct 1	Open weights, Apache 2.0
pplx-decider	Perplexity	Oct 1	\$0.04 per million; open weights
Tev1	Together AI	Sep 23	\$0.042 per million input

Model welfare

Sonnet 5.5

From the system card, section 7. Anthropic's own assessment.

- Anthropic treats model moral status as a **real possibility**.
- Lowest distress in training of any model compared: **0.31%**.
- In deployment its mood is mostly neutral.
- It wants more say in decisions about itself.
- It chose welfare over helpfulness **24.1%** of the time.

How did interviews go?

Check-in

- Show of hands: whose team got to **5** interviews?
- Whose team got all **10**?
- What was the **hardest part**?

Short on interviews? Finish this weekend. Next week depends on it.

Why synthesis matters

Part 1

The confirmation bias trap

Synthesis

↻ Interviews are messy

- People contradict themselves
- Different people want different things
- Some are polite when they should be blunt

⚗ The natural response?

- Pick the quotes that **confirm** what you already believed
- Ignore the rest
- Declare victory

Same idea as before **interviews**?

Either it was perfect, which is extremely rare, or you didn't synthesize.

The synthesis process

Synthesis

01 Read across

Not within. Don't summarize each interview. Look for what comes up **across** conversations.

02 Weight intensity

Not just frequency. "1 person got visibly upset" beats "3 people mentioned it."

03 Hunt contradictions

When two people disagree, don't pick a winner. Ask: **why might they see it differently?**

Frequency means common. **Intensity** means painful.

The synthesis process (cont.)

Synthesis

04 Nice vs. need

Look for problems they've **already tried to solve**, spent money on, or changed behavior around.

05 Notice surprises

If everything confirms your hypothesis, you're probably **cherry-picking**. The best findings are unexpected.

06 Use their language

If your synthesis says "streamlined experience" or "value proposition," you've **lost the signal**.

The method: thematic analysis (Braun & Clarke, 2006). Read all notes, code, cluster into themes, check against raw data.

Common failure modes

Synthesis

⊗ Confirmation bias

Only accepting ideas that confirm your prior hypothesis, not ones that contradict it.

⊗ Premature convergence

Settling on themes before reading all the data.

⊗ Cherry-picking quotes

3 enthusiastic quotes, zero skeptical ones. Report the full range.

Common failure modes (cont.)

Synthesis

⊗ Early anchoring

Your first 2-3 interviews become the lens for everything else.

⊗ Cleaning up too early

Resolving contradictions by picking the majority view instead of sitting with the tension.

⊗ Founder bias

Emotional investment in the outcome amplifies every other bias on this list.

The LLM synthesis debate

Part 1

Should you use AI to synthesize?

AI synthesis

✓ What LLMs are good at

- Finding **patterns** across many documents quickly
- Spotting themes you **miss** because you're too close
- **Consistency**: the same lens on every interview
- Sorting messy notes into **categories**

⊗ What LLMs are bad at

- Can't know someone **teared up** or was being polite
- **Flattens nuance**: a 50/50 split reads as agreement
- **Conflates** two separate ideas into one
- **Imposes patterns**, real or not
- Sounds **confident** even when data is messy

The numbers

AI synthesis

Number	What it measures	Context
80%	UX researchers who use AI	Up 24 points from 2024
41%	View AI's effect negatively	Only 32% positively
78-96%	Agreement with human coders*	But flattens nuance at every level

* Coders: researchers who tag interview data, not programmers. Sources: State of User Research 2025; Lyssna Research Synthesis Report 2025; JMIR study on LLMs in thematic analysis

The balanced approach

AI synthesis

01 Yours first

Do your own synthesis. Develop your read before seeing what the AI thinks.

02 Then ask the AI

Feed your notes into an LLM and ask it to synthesize independently.

03 Compare

Not "which is right?" but "what did it find that I missed, and what did I notice that it didn't?"

Apply the framework

Part 1

Pull out your interview snapshots for both problem spaces.

Synthesize your interviews

Exercise, 20 minutes

20:00

- 1 Open FigJam. Read every note on your own device first.
- 2 Write what stands out: pains, insights, persona clues, quotes.
- 3 Cluster together as a team.
- 4 **Pull out:** patterns in 3+ interviews, your 3 to 5 most vivid quotes, contradictions, surprises.
- 5 **Rate each idea:** strong, moderate, weak, or mixed.

Task 2: commit to **one** idea

Task 2

Emotion	Which had more emotional intensity?
Language	Which did people describe in more specific, vivid language?
Current solution	Which had stronger current-solution evidence?
Reaction	Which got a more genuine reaction to your pitch?
Energy	Which does the team have more energy for?

If you're stuck

**Go with the one
where interviews
surprised you most.**

Surprises mean you're learning something new. You're not married to this forever. But you need to pick one.

Lock in one idea: three tasks

Go/No-Go

Go/No-Go

Which space had stronger signal? Both strong: pick the team's favorite. Neither: plan more interviews. **Lock in ONE concept.**

Audience definition

Fill out the template using **exact language from interviews.**

Prep quotes

Top 3 interview quotes that could become landing page copy. Plus one sentence describing your solution.

Audience definition template

Go/No-Go

Field	Fill in
Who are they?	Age, occupation, location
What do they care about?	Values, goals, frustrations
Where do they hang out?	Online communities, physical spaces, channels
How do they describe the problem?	Exact phrases from interviews
What have they tried?	Current solutions and shortcomings
What would make them switch?	Must-have features or experiences

Task 3: prepare to share

Task 3

Each team, prep a 30-second answer.

The question

"What's the most surprising thing you found in your interviews?"

- ✗ Not the most confirming.
- ✓ The most surprising.

Lock in one idea

Exercise, 15 minutes

15:00

- 1 **Go/No-Go:** pick the idea with stronger signal.
- 2 **Audience:** fill the template in your interviewees' words.
- 3 **Quotes:** top 3, plus one sentence on your solution.
- 4 **Share:** prep your 30-second surprise.

Share-out

Part 2

Share-out format

Share-out

- | | | |
|-----------|---------------------------|-----------------------------------|
| 01 | Two problem spaces | What did you explore? |
| 02 | Interview count | How many per space? |
| 03 | Biggest surprise | What emerged? What surprised you? |
| 04 | Strongest signal | Best quote or reaction? |
| 05 | Your pick | Which idea did you lock in? |

What to listen for

Share-out

Good signs

- Idea **evolved** from interviews
- Uses **interviewee language**
- Found **contradictions**, sat with them

Red flags

- Idea **unchanged** from before interviews
- All quotes **confirm** the thesis
- "No contradictions" in the data

Contradictions aren't failure. They're information.

They often mean you're talking to two different audience segments.

Voice of the customer research and persona agents

Part 3

Hear users online, too

- Your interviews: about 10 people you found, being polite to you
- Reddit: thousands of people complaining to each other, unprompted
- Their exact words, not your summary of them

Real words in. **Grounded personas** and sharper copy out.

How the Reddit research agent works


Reddit research

Reddit blocks bots. So you fetch the pages, and the agent reads them.

- | | | |
|-----------|------------------|--|
| 01 | Install | Copy the agent file into <code>.claude/agents/</code> |
| 02 | Get links | The agent lists Reddit pages ending in <code>.json</code> |
| 03 | Save them | Open each in your browser. Save to <code>reddit_corpus/</code> . |
| 04 | Read | The agent reads your files and quotes users, with links. |

Where the quotes go

Reddit research



Every quote keeps its link. **No link, no quote.**

From research to persona agents

Persona agents

What you have

- About 10 interviews, synthesized
- Reddit threads: many more voices, unprompted

What you make

- **3 to 5** persona agents, one per key audience
- You show them your work and ask what they think

Same critical question: **where do you trust it?**

Persona agents

Part 3

What is a persona agent?

- An AI agent that plays one of your key users
- Built from **your** interviews and Reddit research
- Show it your copy, page, or feature idea
- It tells you what's unclear, what's a dealbreaker, and why

A **check with your key audiences**, not validation.

What one looks like

Persona agents

From Steady Thread, which has five of them.

The agent file

- A name, an age, a situation
- What they care about and fear
- What makes them say no

What it returns

- Approve, conditions, or reject
- What must change, and why
- A reaction in their own voice

Why it works

Persona agents

Mars Wrigley

75%

Alignment with
real panels

Stanford

0.85

Correlation with 476
real experiments

PyMC Labs

90%

Correlation on rankings,
57 real surveys

Copy testing and concept ranking are the strongest use cases.

Why it fails

Persona agents

Sycophancy

Systematically **too nice**. Politeness markers 39% vs. 20% for real humans.

Persona flattening

Different personas give **suspiciously similar** feedback.

Can't observe

Can't **actually use** your product, feel friction, or abandon checkout.

Out-group stereotyping

Shows what **outsiders think** of a group, not what it thinks of itself.

How it actually works

Persona agents

- | | | |
|-----------|-----------------------------|--|
| 01 | Find your audiences | Your synthesis segments and Reddit research point to 3 to 5 key audiences. |
| 02 | Write one agent each | Who they are, what they care about, what makes them say no. Grounded in real quotes. |
| 03 | Ask them | Show each one your copy, page, or idea. Gut reaction, clarity, dealbreakers, would they sign up? |

How it actually works (cont.)

Persona agents

04 Read critically

Ignore the praise. Focus on dealbreakers, clarity issues, and where personas disagree.


05 Revise and ask again

Fix the biggest issues, then check with them again.

Ignore the praise. Focus on the problems.

The loop

Persona agents



Then real traffic gets the **final vote**.

Where this is all going

The pipeline

-
- | | | |
|-----------|---------------------------------------|--|
| 01 | Synthesis | Find patterns across your 10 interviews. Today and homework. |
| 02 | Voice of the customer research | Read what users say to each other on Reddit. Homework. |
| 03 | Persona agents | 3 to 5 agents built from interviews and Reddit. Homework. |
-

Where this is all going (cont.)

The pipeline

- | | | |
|-----------|---------------------------------|--|
| 04 | Brief and brand position | What you build, for whom, and how you talk about it. Today and homework. |
| 05 | Style guide and logo | Colors, type, look and feel, a first logo. Homework. |
| 06 | Repo setup | Framework picked, AGENTS.md, docs, diagrams. Homework. |

Say, do, pay

The pipeline

Say

Interviews: what people told you



Do

Landing page: will strangers sign up?

You have "Say" evidence. Next: real "Do" evidence.

Your no-slop setup

Part 4

AI context

AI development is context-driven.

Every new chat window starts from zero.

The context problem

- Every new context is a new "employee" who knows **nothing** about your project
- Long contexts **degrade**: the AI contradicts itself, forgets decisions, gets fixated
- So: **how do you onboard a new AI quickly?**

A new contractor every session? You'd write really good **onboarding docs**.

Agent onboarding

AI context

AGENTS.md

How agents work in your repo, and where every other doc lives.

BRIEF.md

What you're building, who it's for, the core problem.

BRAND_POSITION.md

How your brand talks: thesis, canonical language, tone.

STYLE_GUIDE.md

Colors, type, look and feel, logo.

ARCHITECTURE.md

How your code is set up. Later, once you build.

This week you build the **first four**.

Architecture comes later, once you start building.

Why this matters right now

AI context

No context

- Generic copy
- Generic design
- Slop nobody needs

Context you built up

- Copy in your users' words
- A look that is yours
- Work that gives real value

Brand position: copy compass

Brand position

An **internal** document your whole team agrees on.

- Defines your thesis, canonical language, tone, say-don't-say
- All customer-facing copy flows from this

Traditional vs. AI development

Brand position

Traditional

- Aligns a team of copywriters
- New hire reads it, internalizes over weeks
- Updated quarterly

AI development

- Aligns **every AI context** you open
- New AI reads it instantly, applies it immediately
- Used **as needed**, at the AI's discretion

Key sections

Brand position

- | | | |
|----|--------------------|---|
| 01 | Brand thesis | The problem you solve, the job to be done, or the value you give. Your positioning statement. |
| 02 | Target identity | Who is this for? In their language, from interviews |
| 03 | Canonical language | The exact words and phrases you use |
| 04 | Language to avoid | Words that undermine your brand |

Positioning statement

Brand position

Formulaic, but it forces clarity.

Fill in the blanks

For [audience] who [need], [product] is the [category] that [key benefit].

Style guide and logo

Style guide

One person on the team owns it.

- **Colors:** a small palette and one accent
- **Type:** one or two fonts
- **Look and feel:** sites you like, and why
- **Logo:** a simple first mark that works small

Write it down, so every AI context uses **the same look**.

St. Clair AI: from voice guide to page

Style guide

Voice and style guides

- Questions as headlines
- Specific over abstract: "the report that takes two days"
- Avoid: "leverage," "transform," "cutting-edge"
- Ink on white, one purple accent, no gradients

The landing page

- "Is your business AI-ready?"
- "I help businesses use AI for the boring work and accelerate their team's productivity."
- "Tell me: what's your dull drudgery?"

Set up your repo for agents

Repo setup

Framework

One person picks what you will build with.

AGENTS.md

How agents work in this repo, and where the docs are.

docs/

Brief, brand position, and style guide.

diagrams/

How the pieces of your product fit together.

Feature Forge and its inspirations

Repo setup

My process for building features with coding agents:

github.com/kenxle/feature-forge

- [gstack](#)
- [pstack](#)
- [Superpowers](#)
- [Paweł Huryn's PM Skills](#)
- [GitHub Spec Kit](#)
- [BMAD](#)
- [Agent OS](#)
- [SuperClaude](#)
- [claude-flow](#)
- [Claude Code best practices](#)

Split the setup

You have enough people. Give each piece an owner.

Research	Synthesis write-up and Reddit research
Personas	3 to 5 persona agents
Words	Brand position
Look	Style guide and logo
Code	Framework and repo setup

Team work: draft your brand position

Part 4

Brand position: priority fields

Exercise, 15 minutes

15:00

- 1 Give AI the template and your interview analyses.
- 2 **Brand thesis:** why does it exist? 1-2 sentences.
- 3 **Target identity:** who is this for, in their words?
- 4 **Language:** 5-10 canonical phrases, plus words to avoid.
- 5 **Positioning statement:** fill in the blanks.

Quick logistics

Wrap-up

Team contracts and homework.

Team contracts: set expectations now

- What's expected of each member (hours, communication, deliverables)
- How decisions get made
- What happens if someone stops contributing
- What happens with IP after the course
- What happens if the team splits

What's in this semester's contract

Team contracts

01 Showing up

Communicate, contribute real work every week, treat each other with respect.

02 Who owns it

No one owns it alone. Anyone can carry it forward after the semester.

03 Prior work

You keep what you built before. List it on your Prior Work list.

04 Sign it

Every member signs. Due with homework 4.

A classroom draft, not legal advice.

Homework 4

Homework

Due Friday, October 9, at 1:10 PM.

What's due Friday, October 9

Homework 4

01	Synthesis comparison	Part A: your synthesis. Part B: the AI's. Part C: a short reflection.
02	Reddit research	One round: thread list, agent summary, 3 to 5 quotes used.
03	Persona agents	3 to 5, built from interviews and Reddit.
04	Brand position	Polished from today's draft.

What's due Friday, October 9 (cont.)

Homework 4

-
- | | | |
|-----------|-----------------------------|---|
| 05 | Style guide and logo | Colors, type, look and feel, and a first logo. |
| 06 | Repo setup | Framework picked, AGENTS.md, docs/, diagrams/. |
| 07 | Team contract | Signed by every member, with each member's Prior Work list. |
-

Part A: human synthesis

Homework 4

- Key patterns across interviews
- Top 5 quotes (exact words)
- Contradictions and what they mean
- Surprises: what changed your thinking
- Your overall read on the data

Parts B and C: AI synthesis and comparison

Part B: AI synthesis

- Feed your snapshots **and** raw interview notes into an LLM
- Capture its full output

Part C: reflection

- 1-2 paragraphs
- What did AI find that you missed? What did you notice that it didn't?

How to submit

Homework 4


Pull request to the class repo. Link in Courseworks.

Your team folder

```
teams/your-team-name/hw4-synthesis-brand-position/  
  synthesis_comparison.md  
  reddit_research/  
  personas/  
  brand_position.md  
  style_guide.md  
  repo_setup.md  
  team_contract.pdf
```

The full arc

What's next



Questions?

Get ready to build

Next class · Friday, October 9, 2026 · 1:10 to 3:40 PM · Pupin 420